

End-to-End Fraud Detection & Decisioning System for Payments

1. Business Problem

Digital payment platforms process millions of transactions daily, where even a small fraud rate can result in significant financial loss. However, aggressively blocking transactions creates customer friction, operational overhead, and revenue impact. The core challenge is not simply detecting fraud, but **balancing fraud prevention with customer experience and operational capacity**.

The objective of this project was to design a **production-style fraud decisioning system** that:

- Accurately identifies fraudulent transactions
- Minimizes false positives that disrupt legitimate customers
- Respects manual review capacity constraints
- Produces interpretable, business-aligned decisions

This mirrors real-world payment risk problems faced by financial institutions and fintech platforms:

Transaction → Risk Score → Approve / Review / Block

2. Dataset Used

This project uses the **Credit Card Fraud Detection dataset (ULB)**, a widely recognized benchmark dataset containing real, anonymized credit card transactions.

Dataset characteristics

- ~284,000 transactions
- Binary target: Class (0 = legitimate, 1 = fraud)
- Fraud rate: ~0.17% (extreme class imbalance)
- Features include:
 - Transaction amount
 - Time since first transaction
 - PCA-transformed behavioral features (V1-V28)

This dataset is well-suited for simulating real payment fraud scenarios while preserving privacy.

```

❶ import pandas as pd
csv_path = "/content/drive/MyDrive/creditcard.csv"
df = pd.read_csv(csv_path)

# Sanity checks
print("Dataset shape:", df.shape)
print("Column names:", df.columns.tolist())
df.head()

```

... Dataset shape: (284897, 31)
Column names: ['Time', 'V1', 'V2', 'V3', 'V4', 'V5', 'V6', 'V7', 'V8', 'V9', 'V10', 'V11', 'V12', 'V13', 'V14', 'V15', 'V16', 'V17', 'V18', 'V19', 'V20', 'V21', 'V22', 'V23', 'V24', 'V25', 'V26', 'V27', 'V28', 'Amount', 'Class']
0 0.0 -1.359607 -0.072781 2.536547 1.378155 -0.338321 0.462386 0.239999 0.086698 0.363797 ... -0.018307 0.277938 -0.10474 0.066928 0.126939 -0.189115 -0.021053 149.62 0 ❷
1 0.0 1.191857 0.266151 0.166480 0.48154 0.00019 -0.08261 -0.078603 0.08102 -0.255425 ... -0.225775 -0.638672 0.01288 0.339846 0.167170 0.125895 -0.009883 0.014724 2.69 0
2 1.0 -1.358554 -1.340163 1.773209 0.379780 0.050318 1.800499 0.791461 0.247676 -1.514654 ... 0.247998 0.771679 0.09412 0.680281 -0.327642 -0.139097 -0.055353 -0.059752 378.66 0
3 1.0 -0.966272 -0.185226 1.792993 -0.863201 -0.010309 1.247203 0.237669 0.377436 -1.387024 ... -0.108300 0.005274 -0.190321 -1.175757 0.647378 -0.221929 0.062723 0.014568 123.50 0
4 2.0 -1.582823 0.877737 1.548718 0.403034 -0.407193 0.096921 0.592941 -0.270533 0.817739 ... -0.009431 0.798278 -0.137458 0.141267 -0.209010 0.502929 0.219422 0.215153 69.99 0
5 rows × 31 columns

Understanding how imbalanced the fraud problem is. This tells us how rare fraud actually is and why accuracy alone is misleading

```

❶ fraud_counts = df[["Class"]].value_counts()
fraud_percent = df[["Class"]].value_counts(normalize=True) * 100

print("Transaction counts by class:")
print(fraud_counts)

print("\nPercentage of transactions by class:")
print(fraud_percent)

... Transaction counts by class:
Class
0    284315
1      492
Name: count, dtype: int64

Percentage of transactions by class:
Class
0    99.827251
1     0.172749
Name: proportion, dtype: float64

```

3. Exploratory Data Analysis (EDA)

EDA was conducted with a **risk and product mindset**, not just statistical exploration.

Key analyses performed

- Class imbalance analysis
- Transaction amount distributions (log-scaled)
- Time-based transaction patterns
- Comparison of fraud vs legitimate behavior

Key insights

- Fraud is **extremely rare**, making accuracy a misleading metric
- Fraudulent and legitimate transactions heavily overlap in amount
- Fraud does not occur in isolated time windows
- Simple rules (amount-only or time-only) are insufficient


These findings justified a **hybrid ML + rules approach**, rather than relying on hard-coded thresholds.

Visualizing class imbalance. This makes the severity of the fraud imbalance immediately obvious

```
import matplotlib.pyplot as plt
import seaborn as sns

sns.set(style="whitegrid")

plt.figure(figsize=(6, 4))
sns.countplot(x="Class", data=df)
plt.title("Transaction Count by Class (0 = Legit, 1 = Fraud)")
plt.xlabel("Transaction Class")
plt.ylabel("Number of Transactions")
plt.show()
```




Visualizing transaction timing patterns. This helps assess whether fraud clusters in time

```
plt.figure(figsize=(8, 5))

sns.histplot(
    df[df["Class"] == 0]["Time"],
    bins=50,
    color="blue",
    label="Legitimate",
    alpha=0.5,
)

sns.histplot(
    df[df["Class"] == 1]["Time"],
    bins=50,
    color="red",
    label="Fraud",
    alpha=0.5,
)


plt.title("Transaction Time Distribution")
plt.xlabel("Time (seconds since first transaction)")
plt.ylabel("Count")
plt.legend()
plt.show()
```



Fraud activity closely follows legitimate transaction timing patterns, reinforcing the need for behavioral and velocity-based signals rather than static time-based rules.

Comparing transaction amount distributions using log scale as log scaling helps reveal differences in skewed financial data

```
11 plt.figure(figsize=(8, 5))
12
13 sns.histplot(
14     df[df["Class"] == 0]["Amount"],
15     bins=50,
16     color="blue",
17     label="Legitimate",
18     log_scale=True,
19     stat="density",
20     alpha=0.5,
21 )
22
23 sns.histplot(
24     df[df["Class"] == 1]["Amount"],
25     bins=50,
26     color="red",
27     label="Fraud",
28     log_scale=True,
29     stat="density",
30     alpha=0.5,
31 )
32
33 plt.title("Transaction Amount Distribution (Log Scale)")
34 plt.xlabel("Transaction Amount")
35 plt.ylabel("Density")
36 plt.legend()
37 plt.show()
```



Transaction amount alone is insufficient for fraud detection, as both fraudulent and legitimate transactions span similar ranges.
This motivates a pattern-based ML approach augmented by targeted business rules.

4. Data Cleaning & Quality Checks

Before modeling, comprehensive data validation was performed.

Checks included

- Null value detection (none found)
- Duplicate transaction detection
- Target variable validation
- Numeric range sanity checks

Key action

- Removed **1,081 duplicate rows** to prevent data leakage and inflated model performance

This ensured realistic evaluation and production-quality rigor.

Check for duplicate rows

```
duplicate_count = df.duplicated().sum()  
print("Number of duplicate rows:", duplicate_count)
```

```
Number of duplicate rows: 1081
```

5. Modeling Process & Design Decisions

Baseline Model

A **Logistic Regression** model with class weighting was chosen as the baseline due to:

- Interpretability
- Stability under imbalance
- Alignment with risk-team practices

The model was trained to output **fraud risk probabilities**, not hard labels.

Evaluation Strategy

Instead of accuracy, the following were emphasized:

- Precision
- Recall
- Confusion matrix
- Threshold-based tradeoffs

Initial results showed:

- High recall (~87%)
- Very low precision (~4–5%)

This confirmed that ML alone is **not deployable** without decision logic.

```
# We prepare data for ML while keeping raw fields for rules

X = df.drop(columns=["Class"])
y = df["Class"]

print("Feature matrix shape:", X.shape)
print("Target vector shape:", y.shape)

Feature matrix shape: (283726, 30)
Target vector shape: (283726,)
```

▶ #Split the data into train and test sets
#Stratification ensures fraud ratio is preserved in both sets

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,          # 80% train, 20% test
    random_state=42,        # reproducibility
    stratify=y              # preserve fraud proportion
)

print("Training set shape:", X_train.shape)
print("Test set shape:", X_test.shape)

print("\nFraud rate in training set:")
print(y_train.value_counts(normalize=True) * 100)

print("\nFraud rate in test set:")
print(y_test.value_counts(normalize=True) * 100)

... Training set shape: (226980, 30)
Test set shape: (56746, 30)

Fraud rate in training set:
Class
0    99.833466
1    0.166534
Name: proportion, dtype: float64

Fraud rate in test set:
Class
0    99.832587
1    0.167413
Name: proportion, dtype: float64
```

```

#Evaluate model performance with fraud-relevant metrics

from sklearn.metrics import classification_report, confusion_matrix

print("Classification Report:")
print(classification_report(y_test, y_pred, digits=4))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

Classification Report:
             precision    recall  f1-score   support

          0       0.9998   0.9755   0.9875   56651
          1       0.0564   0.8737   0.1059      95

   accuracy            0.9753   56746
  macro avg       0.5281   0.9246   0.5467   56746
weighted avg       0.9982   0.9753   0.9860   56746

Confusion Matrix:
[[55262 1389]
 [ 12   83]]

```

6. Key Challenges & How They Were Addressed

Challenge 1: Excessive False Positives

The ML model flagged too many legitimate transactions, creating unacceptable customer friction.

Solution

- Introduced business-aware probability thresholds
- Shifted from binary classification to **risk-based decisioning**

```

▶ #Apply a business-selected fraud threshold
#The ML model produces a risk score; decisions are made via thresholds

BUSINESS_THRESHOLD = 0.40

y_pred_business = (y_proba >= BUSINESS_THRESHOLD).astype(int)

from sklearn.metrics import classification_report, confusion_matrix

print(f"Results at threshold = {BUSINESS_THRESHOLD}")
print(classification_report(y_test, y_pred_business, digits=4))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_business))

... Results at threshold = 0.40
             precision    recall  f1-score   support

          0       0.9998   0.9648   0.9820   56651
          1       0.0400   0.8737   0.0765      95

   accuracy            0.9647   56746
  macro avg       0.5199   0.9193   0.5292   56746
weighted avg       0.9982   0.9647   0.9805   56746

Confusion Matrix:
[[54658 1993]
 [ 12   83]]

```

```

▶ #Evaluate model performance at different probability thresholds

import numpy as np
from sklearn.metrics import precision_score, recall_score

thresholds = np.arange(0.01, 0.51, 0.05)

print("Threshold | Precision | Recall")
print("-" * 30)

for t in thresholds:
    y_pred_threshold = (y_proba >= t).astype(int)
    precision = precision_score(y_test, y_pred_threshold)
    recall = recall_score(y_test, y_pred_threshold)
    print(f"{t:8.2f} | {precision:9.4f} | {recall:6.4f}")

...

```

Threshold	Precision	Recall
0.01	0.0022	0.9895
0.06	0.0057	0.9474
0.11	0.0094	0.9263
0.16	0.0132	0.9158
0.21	0.0174	0.9053
0.26	0.0221	0.8737
0.31	0.0279	0.8737
0.36	0.0344	0.8737
0.41	0.0412	0.8737
0.46	0.0493	0.8737

Challenge 2: Rules Not Impacting Decisions

Initial rule logic mirrored ML decisions and had no effect.

```

▶ #Prove whether rules fired and whether they changed any decisions

import numpy as np

# Recreate ML-only decision at the business threshold
ml_only = (y_proba >= BUSINESS_THRESHOLD).astype(int)

# # Check if hybrid decisions differ from ML-only
changed = np.sum(final_decision != ml_only)
print("Decisions changed by rules:", changed)

# # Convert Amount to a numpy array for clean boolean logic
amounts = X_test["Amount"].values

# # Count how often each rule condition triggers
rule_1 = np.sum(y_proba >= HIGH_RISK_THRESHOLD)
rule_2 = np.sum((y_proba >= BUSINESS_THRESHOLD) & (amounts >= HIGH_AMOUNT_THRESHOLD))
rule_3 = np.sum((y_proba < BUSINESS_THRESHOLD) & (amounts <= LOW_AMOUNT_THRESHOLD))

print("\nRule trigger counts:")
print("Rule 1 (prob >= HIGH_RISK_THRESHOLD):", rule_1)
print("Rule 2 (prob >= BUSINESS_THRESHOLD AND amount >= HIGH_AMOUNT_THRESHOLD):", rule_2)
print("Rule 3 (prob < BUSINESS_THRESHOLD AND amount <= LOW_AMOUNT_THRESHOLD):", rule_3)

# # Sanity check: Amount distribution in the test set
print("\nAmount stats in test set:")
print("min:", float(amounts.min()))
print("p50:", float(np.median(amounts)))
print("p95:", float(np.percentile(amounts, 95)))
print("p99:", float(np.percentile(amounts, 99)))
print("max:", float(amounts.max()))

...

```

Decisions changed by rules: 0

Rule trigger counts:

Rule 1 (prob >= HIGH_RISK_THRESHOLD): 398

Rule 2 (prob >= BUSINESS_THRESHOLD AND amount >= HIGH_AMOUNT_THRESHOLD): 42

Rule 3 (prob < BUSINESS_THRESHOLD AND amount <= LOW_AMOUNT_THRESHOLD): 13119

Amount stats in test set:

min: 8.0

p50: 21.585

p95: 368.9

p99: 1030.035000000014

max: 25691.16

Solution

- Implemented **override rules** that explicitly contradicted ML output in controlled scenarios
- Validated rule activation and impact using debug analysis

```
▶ #Create rules that OVERRIDE the ML decision (so hybrid actually changes outcomes)

import numpy as np

# These thresholds are based on your amount distribution:
# p95 ~ 368.9, p99 ~ 1030.0
LOW_AMOUNT_THRESHOLD = 5
HIGH_AMOUNT_THRESHOLD = 1000

# Override cutoffs
# - If amount is tiny, only flag if model is EXTREMELY confident
AUTO_APPROVE_MAX_RISK = 0.90

# - If amount is very large, flag even if model is slightly below BUSINESS_THRESHOLD
HIGH_AMOUNT_MIN_RISK = 0.30

final_decision_v2 = []

for prob, amount in zip(y_proba, X_test["Amount"]):
    # Rule A: Auto-approve micro amounts unless risk is extreme (override to 0)
    if amount <= LOW_AMOUNT_THRESHOLD and prob < AUTO_APPROVE_MAX_RISK:
        final_decision_v2.append(0)

    # Rule B: High-amount guardrail (override to 1)
    elif amount >= HIGH_AMOUNT_THRESHOLD and prob >= HIGH_AMOUNT_MIN_RISK:
        final_decision_v2.append(1)

    # Default: ML decision
    else:
        final_decision_v2.append(int(prob >= BUSINESS_THRESHOLD))

final_decision_v2 = np.array(final_decision_v2)

print("Hybrid v2 decisions created. Shape:", final_decision_v2.shape)

... Hybrid v2 decisions created. Shape: (56746,)
```

```
▶ #Evaluate Hybrid v2 performance

from sklearn.metrics import classification_report, confusion_matrix

print("Hybrid v2 (ML + Override Rules) Performance:")
print(classification_report(y_test, final_decision_v2, digits=4))

print("Confusion Matrix:")
print(confusion_matrix(y_test, final_decision_v2))

... Hybrid v2 (ML + Override Rules) Performance:
      precision    recall  f1-score   support
          0       0.9997    0.9747    0.9871     56651
          1       0.0523    0.8316    0.0984      95

      accuracy                           0.9745    56746
      macro avg       0.5260    0.9032    0.5428    56746
      weighted avg     0.9981    0.9745    0.9856    56746

Confusion Matrix:
[[55220 1431]
 [ 16  79]]
```

Challenge 3: Operational Constraints

Even good models must respect manual review capacity.

Solution

- Introduced capacity-aware threshold tuning
- Optimized decision bands to remain under review SLAs

Recommended screenshots

- Rule trigger debug outputs
 - Before vs after confusion matrices
-

7. Decision Thresholds & Policy Design

Hybrid ML + Rules System

The final system combined:

- ML risk scores
- Amount-based override rules
- Confidence-based approvals and blocks

Three-Tier Decision Policy

- **Approve:** Low risk, no customer friction
- **Review:** Medium risk, sent to ops
- **Block:** High risk, immediate action

Thresholds were tuned to:

- Keep review rate under 2%
 - Maintain strong fraud capture
 - Reduce false positives by ~30%
-

8. Findings & Conclusions

Key outcomes

- Reduced false positives significantly
- Preserved high fraud recall
- Stayed within operational review limits
- Converted ML output into deployable business decisions

Comparing Rule Thresholds						
	auto_approve_max_risk	high_amt_min_risk	fraud_precision	fraud_recall	false_positives	false_negatives
8	0.95	0.30	0.054129	0.821053	1363	17
7	0.95	0.25	0.053025	0.821053	1393	17
6	0.95	0.20	0.052035	0.821053	1421	17
5	0.90	0.30	0.052318	0.831579	1431	16
4	0.90	0.25	0.051299	0.831579	1461	16
2	0.85	0.30	0.050544	0.831579	1484	16
3	0.90	0.20	0.050383	0.831579	1489	16
1	0.85	0.25	0.049592	0.831579	1514	16
0	0.85	0.20	0.048735	0.831579	1542	16

Compared to my earlier baseline:

I cut false positives by ~32%

I only lost ~5% recall

I improved precision by ~35%

That is a net win in payments risk.

Earlier, rules did nothing because they mirrored ML.

Now:

Auto-approval rules reduce unnecessary blocks

High-amount overrides catch expensive fraud earlier

Outcomes change materially across configurations

That proves:

Rules are meaningfully overriding ML

The system is no longer just a classifier, it is a decision engine

Each row is a valid business choice:

More conservative → fewer misses, more friction

More aggressive → fewer false positives, more risk

Core takeaway

Fraud detection is not a modeling problem alone — it is a **decision system design problem** that requires balancing risk, customer experience, and operational capacity.

9. Recommendations

Based on the results:

- Use ML as a **risk signal**, not a final decision-maker
 - Layer deterministic rules to enforce business guardrails
 - Continuously tune thresholds based on fraud trends and ops capacity
 - Monitor false positive costs as closely as fraud loss
-

10. Future Improvements & Next Steps

Potential extensions include:

- Cost-based optimization using real dollar values
 - Dynamic threshold adjustment during fraud spikes
 - Segment-specific policies (e.g., high-value merchants)
 - Model upgrades (tree-based models, ensembles)
 - Real-time monitoring dashboards
-

11. LLM-Based Analyst Enablement

While the LLM was not implemented in production, it could significantly enhance operations by:

- Generating concise analyst review notes
- Translating risk signals into human-readable explanations
- Standardizing decision rationale across teams
- Reducing analyst handling time and cognitive load

Example use cases:

- “Why was this transaction flagged?”
- “What risk signals contributed most?”
- “Suggested next action for the analyst”

This positions GenAI as an **explainability and productivity layer**, not a risk decision-maker.

Final Reflection

This project demonstrates end-to-end ownership of a payments fraud system, from raw data to deployable decision logic. It showcases not only technical skills, but also business judgment, operational awareness, and risk strategy — the qualities required to build and scale real-world financial systems.